

# Election Violence Newspaper Article Selection Process Pilot: Evaluate Keyword Selection Method

August 9, 2018

## 1 Summary

After the code to setup and download data (Sections 2-??) the document below includes: First, an implementation of code to select of articles for coding, where initial evaluation of articles is based on keywords from the boolean searches which return those articles (Sections ??-??). This classification is done using actual searches conducted (effectively 4 keywords) with a Naive Bayes classifier.

Second, the document then provides an evaluation of the keyword selection step based on user classified articles from the platform with Naive Bayes and other classifiers and with 4 and 10 keyword sets (Sections 4-7).

Finally, to provide an alternative way of evaluating the process, the document includes 10 randomly selected cases which are rejected for coding in the implementation code (Section ??).

## 2 Packages used

```
library(reportRx) # note using this just for sanitizing LaTeX output, not included in durhamevp package
library(tidyverse)
library(quanteda)
library(durhamevp)
library(RTextTools) # not sure if this is included as part of the durhamevp package
```

## 3 Download Classified Cases

There is an initial set of cases which represent the concept (in this case ‘election violence’). We also have a set of cases which do not represent the concept (Note: King et al do not have the non-election violence cases).

I also add indicators that this material is classified (`classified==1`) and that it is not unclassified (`unclass==0`) which will be useful when added to subsequent material.

```
# In fact we have some classified docs non-EV cases here King's R set doesn't contain any non-cases
# I think we should use our more extensive information here

classdocs<-durhamevp::get_classified_docs()

classdocs<-classdocs %>%
  dplyr::mutate(std_url = sub("download/", "", url)) %>%
  dplyr::mutate(unclass=0, classified=1)
```

## 4 Assess Naive Bayes selection performance on 4 keywords (binary dfm)

Here there are four keywords (election, riot, disturbance, incident) which are either present (1) or not-present (0). Use these patterns to identify election violence articles with a Naive Bayes classifier.

```
class_corpus<-quantda::corpus(classdocs[classdocs$selection_article==1,], text_field="ocr")
keywords<-c("election", "riot", "incident", "disturbance")
class_dfm_4<-quantda::dfm(class_corpus, select=keywords)
#class_dfm<-preprocess_corpus(class_corpus, stem=FALSE, min_termfreq = 20)

the_sets_4<-split_dfm(class_dfm_4, n_train = 1000)
classifier<-quantda::textmodel_nb(the_sets_4$training_set, y=quantda::docvars(the_sets_4$training_set)
quantda::docvars(the_sets_4$testing_set, "predicted")<-predict(classifier, newdata = the_sets_4$testing_set)
caret::confusionMatrix(data=quantda::docvars(the_sets_4$testing_set, "predicted"), reference=factor(quantda::docvars(the_sets_4$testing_set)))

## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0  10    8
##           1 134  395
##
##              Accuracy : 0.7404
##              95% CI : (0.7015, 0.7767)
##      No Information Rate : 0.7367
##      P-Value [Acc > NIR] : 0.4451
##
##              Kappa : 0.069
##  Mcnemar's Test P-Value : <2e-16
##
##              Precision : 0.7467
##              Recall : 0.9801
##              F1 : 0.8476
##              Prevalence : 0.7367
##      Detection Rate : 0.7221
##      Detection Prevalence : 0.9671
##      Balanced Accuracy : 0.5248
##
##      'Positive' Class : 1
##
```

## 5 How do other classifiers perform on 4 keywords (binary dfm)

Is performance improved by using other classifiers than Naive Bayes? Here there are four keywords (election, riot, disturbance, incident) which are either present (1) or not-present (0). Use these patterns to identify election violence articles.

```
class_corpus<-quantda::corpus(classdocs[classdocs$selection_article==1,], text_field="ocr")
keywords<-c("election", "riot", "incident", "disturbance")
class_dfm<-quantda::dfm(class_corpus, select=keywords)
class_dfm_bin<-quantda::dfm_weight(class_dfm, scheme="boolean")
doc_matrix<-quantda::convert(class_dfm_bin, "tm")
```

```

## Warning in dfm2tm(x): converted DocumentTermMatrix will not have weight attributes set
correctly

training_nos<-which(1:nrow(class_dfm_bin) %in% sample(1:nrow(class_dfm_bin), 1000))
testing_nos<-which(!1:nrow(class_dfm_bin) %in% training_nos)
container <- RTextTools::create_container(doc_matrix, quanteda::docvars(class_dfm, "EV_article"), train

models_res<-RTextTools::train_models(container, algorithms=c("MAXENT", "SVM", "BOOSTING", "BAGGING", "R
results_res<-RTextTools::classify_models(container, models_res)
analytics_res<-RTextTools::create_analytics(container, results_res)

class_summary<-analytics_res@document_summary %>%
  tibble::rowid_to_column("document_id") %>%
  tidyr::gather(variable, value, -MANUAL_CODE, -document_id) %>%
  separate(variable, c("variable", "var2")) %>%
  filter(var2 %in% c("CODE")) %>%
  unite(value, var2, value) %>%
  group_by(variable, MANUAL_CODE, value) %>%
  tally() %>%
  spread(value, n)

## Warning: attributes are not identical across measure variables;
## they will be dropped
## Warning: package 'bindrcpp' was built under R version 3.4.4

class_summary_by_classifier<-analytics_res@document_summary %>%
  tibble::rowid_to_column("document_id") %>%
  tidyr::gather(variable, value, -MANUAL_CODE, -document_id) %>%
  separate(variable, c("variable", "var2")) %>%
  filter(var2=="LABEL") %>%
  unite(value, var2, value) %>%
  group_by(variable, MANUAL_CODE, value) %>%
  tally() %>%
  spread(value, n)

## Warning: attributes are not identical across measure variables;
## they will be dropped

knitr::knit_print(knitr::kable(class_summary))

```

variable	MANUAL_CODE	CODE_0	CODE_1
CONSENSUS	0	7	177
CONSENSUS	1	6	357
PROBABILITY	0	7	177
PROBABILITY	1	6	357

```
knitr::knit_print(knitr::kable(class_summary_by_classifier))
```

variable	MANUAL_CODE	LABEL_0	LABEL_1
BAGGING	0	7	177
BAGGING	1	6	357
FORESTS	0	NA	184
FORESTS	1	NA	363
LOGITBOOST	0	184	NA
LOGITBOOST	1	121	242
MAXENTROPY	0	7	177
MAXENTROPY	1	6	357
SVM	0	85	99
SVM	1	34	329
TREE	0	7	177
TREE	1	6	357

```

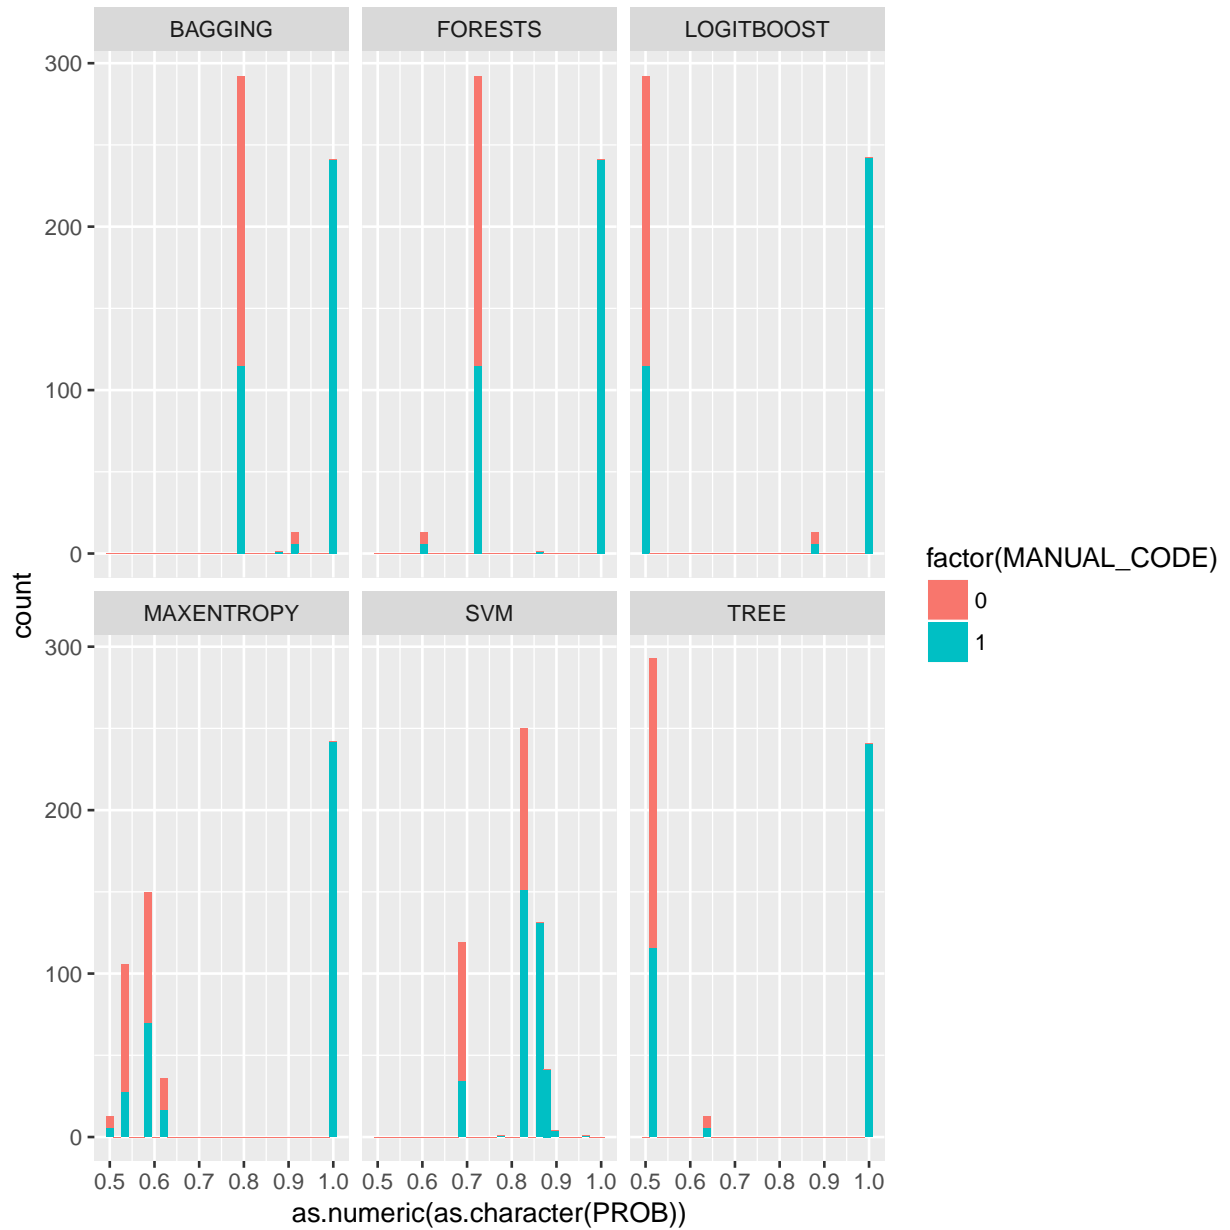
prob_summary<-analytics_res@document_summary %>%
  tibble::rowid_to_column("document_id") %>%
  tidyr::gather(variable, value, -MANUAL_CODE, -document_id) %>%
  separate(variable, c("variable", "var2")) %>%
  filter(var2 %in% c("LABEL", "PROB")) %>%
  spread(var2, value) %>%
  group_by(variable, MANUAL_CODE, LABEL)

## Warning: attributes are not identical across measure variables;
## they will be dropped

prob_summary%>%
  ggplot(aes(x=as.numeric(as.character(PROB)), fill=factor(MANUAL_CODE)))+
  facet_wrap(~variable)+
  geom_histogram()

## 'stat_bin()' using 'bins = 30'. Pick better
## value with 'binwidth'.

```



```
t(analytics_res@algorithm_summary)
```

```
##           0      1
## SVM_PRECISION 0.71 0.77
## SVM_RECALL    0.46 0.91
## SVM_FSCORE    0.56 0.83
## LOGITBOOST_PRECISION 0.60 1.00
## LOGITBOOST_RECALL 1.00 0.67
## LOGITBOOST_FSCORE 0.75 0.80
## BAGGING_PRECISION 0.54 0.67
## BAGGING_RECALL   0.04 0.98
## BAGGING_FSCORE   0.07 0.80
## FORESTS_PRECISION  NaN 0.66
## FORESTS_RECALL   0.00 1.00
```

```

## FORESTS_FSCORE      NaN 0.80
## TREE_PRECISION      0.54 0.67
## TREE_RECALL         0.04 0.98
## TREE_FSCORE         0.07 0.80
## MAXENTROPY_PRECISION 0.54 0.67
## MAXENTROPY_RECALL   0.04 0.98
## MAXENTROPY_FSCORE   0.07 0.80

t(analytics_res@label_summary)

##
##
## NUM_MANUALLY_CODED      184.000000
## NUM_CONSENSUS_CODED     13.000000
## NUM_PROBABILITY_CODED   13.000000
## PCT_CONSENSUS_CODED     7.065217
## PCT_PROBABILITY_CODED   7.065217
## PCT_CORRECTLY_CODED_CONSENSUS 3.804348
## PCT_CORRECTLY_CODED_PROBABILITY 3.804348
##
## 1
## NUM_MANUALLY_CODED      363.00000
## NUM_CONSENSUS_CODED     534.00000
## NUM_PROBABILITY_CODED   534.00000
## PCT_CONSENSUS_CODED     147.10744
## PCT_PROBABILITY_CODED   147.10744
## PCT_CORRECTLY_CODED_CONSENSUS 98.34711
## PCT_CORRECTLY_CODED_PROBABILITY 98.34711

analytics_res@ensemble_summary

##          n-ENSEMBLE COVERAGE n-ENSEMBLE RECALL
## n >= 1          1.00          0.67
## n >= 2          1.00          0.67
## n >= 3          1.00          0.67
## n >= 4          1.00          0.67
## n >= 5          0.81          0.76
## n >= 6          0.44          1.00

summary(analytics_res)

## ENSEMBLE SUMMARY
##
##          n-ENSEMBLE COVERAGE n-ENSEMBLE RECALL
## n >= 1          1.00          0.67
## n >= 2          1.00          0.67
## n >= 3          1.00          0.67
## n >= 4          1.00          0.67
## n >= 5          0.81          0.76
## n >= 6          0.44          1.00
##
##
## ALGORITHM PERFORMANCE
##
##          SVM_PRECISION          SVM_RECALL
##          0.740          0.685
##          SVM_FSCORE LOGITBOOST_PRECISION

```

```
##          0.695          0.800
## LOGITBOOST_RECALL LOGITBOOST_FSCORE
##          0.835          0.775
## BAGGING_PRECISION BAGGING_RECALL
##          0.605          0.510
## BAGGING_FSCORE FORESTS_PRECISION
##          0.435          0.330
## FORESTS_RECALL FORESTS_FSCORE
##          0.500          0.400
## TREE_PRECISION TREE_RECALL
##          0.605          0.510
## TREE_FSCORE MAXENTROPY_PRECISION
##          0.435          0.605
## MAXENTROPY_RECALL MAXENTROPY_FSCORE
##          0.510          0.435
```

## 6 Assess Naive Bayes selection performance on 10 keywords (binary dfm)

Here there are ten keywords ((election, riot, disturbance, incident, mob, stone, window, candidate, party, hustings, magistrate) which are either present (1) or not-present (0). Use these patterns to identify election violence articles with a Naive Bayes classifier.

```
class_corpus<-quantda::corpus(classdocs[classdocs$selection_article==1,], text_field="ocr")
keywords<-c("election", "riot", "incident", "disturbance", "mob", "stone", "window", "candidate", "party", "hustings", "magistrate")
class_dfm_10<-quantda::dfm(class_corpus, select=keywords)
#class_dfm<-preprocess_corpus(class_corpus, stem=FALSE, min_termfreq = 20)

the_sets_10<-split_dfm(class_dfm_10, n_train = 1000)
classifier_10<-quantda::textmodel_nb(the_sets_10$training_set, y=quantda::docvars(the_sets_10$training_set)$selection_article)
quantda::docvars(the_sets_10$testing_set, "predicted")<-predict(classifier_10, newdata = the_sets_10$testing_set)
caret::confusionMatrix(data=quantda::docvars(the_sets_10$testing_set, "predicted"), reference=factor(quantda::docvars(the_sets_10$testing_set)$selection_article, levels=c(0,1)))

## Confusion Matrix and Statistics
##
##          Reference
## Prediction  0    1
##          0  88  52
##          1  73 334
##
##               Accuracy : 0.7715
##               95% CI : (0.734, 0.806)
##       No Information Rate : 0.7057
##       P-Value [Acc > NIR] : 0.0003303
##
##               Kappa : 0.4281
##  Mcnemar's Test P-Value : 0.0736383
##
##               Precision : 0.8206
##               Recall : 0.8653
##               F1 : 0.8424
##               Prevalence : 0.7057
```

```
##          Detection Rate : 0.6106
##    Detection Prevalence : 0.7441
##      Balanced Accuracy : 0.7059
##
##      'Positive' Class : 1
##
```

## 7 How do other classifiers perform on 10 keywords (binary dfm)

The same but with ten keywords (election, riot, disturbance, incident, mob, stone, window, candidate, party, hustings, magistrate) which are either present (1) or not-present (0). Use these patterns to identify election violence articles using a range of other classifiers.

```
class_corpus<-quantda::corpus(classdocs[classdocs$election_article==1,], text_field="ocr")
keywords<-c("election", "riot", "incident", "disturbance", "mob", "stone", "window", "candidate", "party", "hustings", "magistrate")
class_dfm<-quantda::dfm(class_corpus, select=keywords)
class_dfm_bin<-quantda::dfm_weight(class_dfm, scheme="boolean")
doc_matrix<-quantda::convert(class_dfm_bin, "tm")

## Warning in dfm2tm(x): converted DocumentTermMatrix will not have weight attributes set
## correctly

training_nos<-which(1:nrow(class_dfm_bin) %in% sample(1:nrow(class_dfm_bin), 1000))
testing_nos<-which(!1:nrow(class_dfm_bin) %in% training_nos)
container <- RTextTools::create_container(doc_matrix, quantda::docvars(class_dfm, "EV_article"), train=training_nos, test=testing_nos)

models_res<-RTextTools::train_models(container, algorithms=c("MAXENT", "SVM", "BOOSTING", "BAGGING", "RANDOM_FOREST"))
results_res<-RTextTools::classify_models(container, models_res)
analytics_res<-RTextTools::create_analytics(container, results_res)

class_summary<-analytics_res@document_summary %>%
  tibble::rowid_to_column("document_id") %>%
  tidyr::gather(variable, value, -MANUAL_CODE, -document_id) %>%
  separate(variable, c("variable", "var2")) %>%
  filter(var2 %in% c("CODE")) %>%
  unite(value, var2, value) %>%
  group_by(variable, MANUAL_CODE, value) %>%
  tally() %>%
  spread(value, n)

## Warning: attributes are not identical across measure variables;
## they will be dropped

class_summary_by_classifier<-analytics_res@document_summary %>%
  tibble::rowid_to_column("document_id") %>%
  tidyr::gather(variable, value, -MANUAL_CODE, -document_id) %>%
  separate(variable, c("variable", "var2")) %>%
  filter(var2=="LABEL") %>%
  unite(value, var2, value) %>%
  group_by(variable, MANUAL_CODE, value) %>%
  tally() %>%
  spread(value, n)

## Warning: attributes are not identical across measure variables;
## they will be dropped
```



```
knitr::knit_print(knitr::kable(class_summary))
```

variable	MANUAL_CODE	CODE_0	CODE_1
CONSENSUS	0	146	39
CONSENSUS	1	3	359
PROBABILITY	0	146	39
PROBABILITY	1	3	359

```
knitr::knit_print(knitr::kable(class_summary_by_classifier))
```

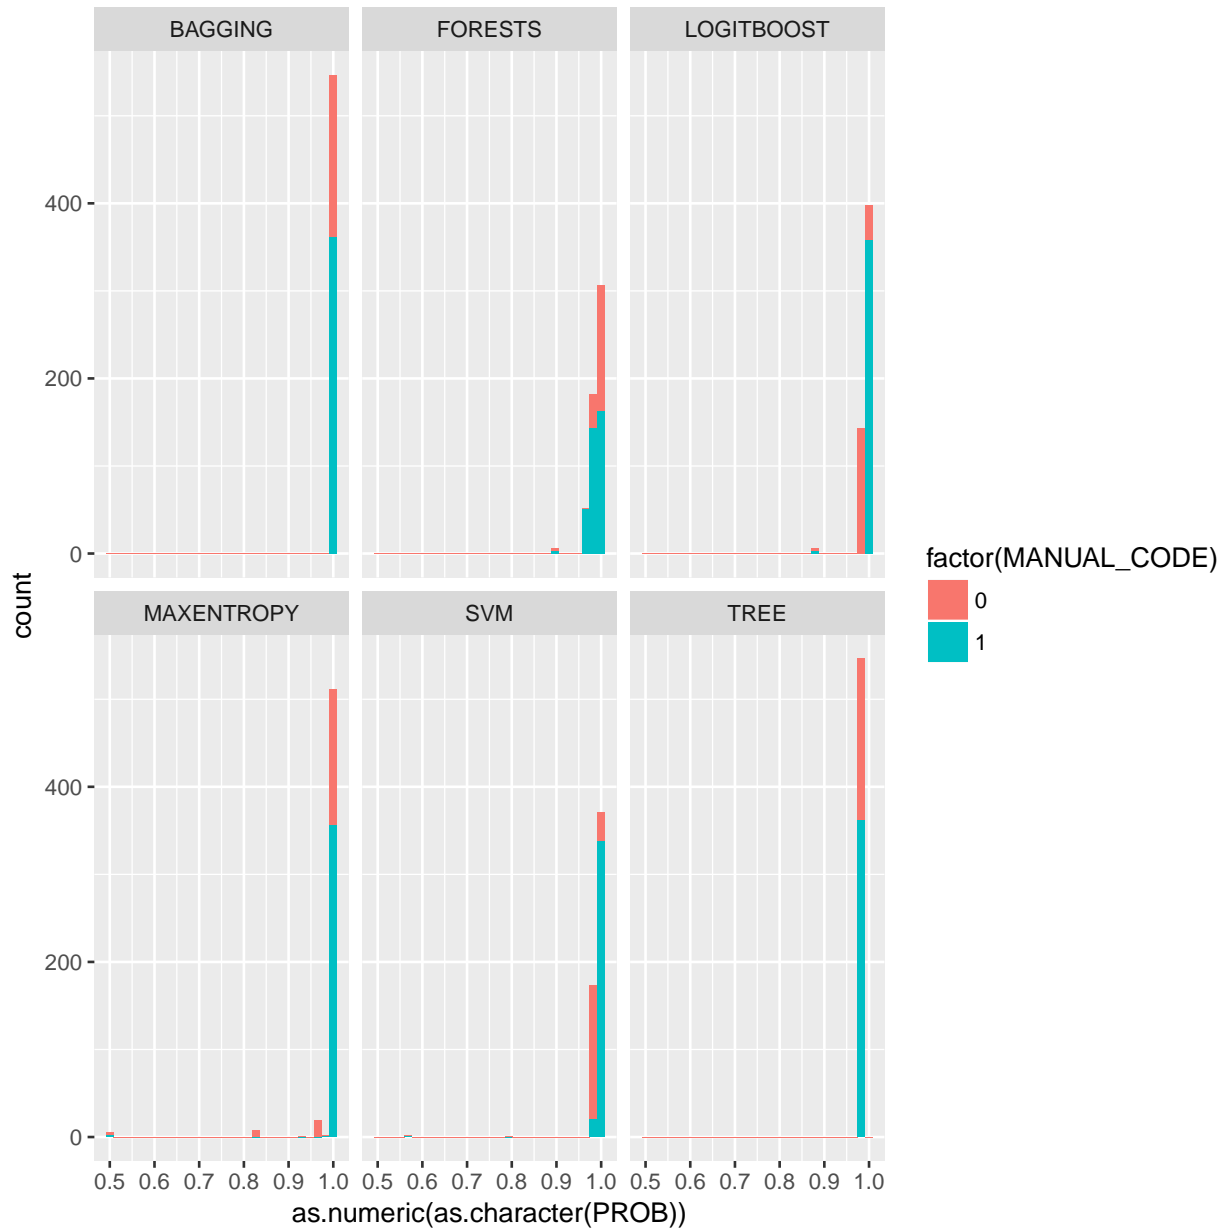
variable	MANUAL_CODE	LABEL_0	LABEL_1
BAGGING	0	146	39
BAGGING	1	3	359
FORESTS	0	143	42
FORESTS	1	NA	362
LOGITBOOST	0	146	39
LOGITBOOST	1	3	359
MAXENTROPY	0	144	41
MAXENTROPY	1	NA	362
SVM	0	146	39
SVM	1	3	359
TREE	0	144	41
TREE	1	NA	362

```
prob_summary<-analytics_res@document_summary %>%
  tibble::rowid_to_column("document_id") %>%
  tidyr::gather(variable, value, -MANUAL_CODE, -document_id) %>%
  separate(variable, c("variable", "var2")) %>%
  filter(var2 %in% c("LABEL", "PROB")) %>%
  spread(var2, value) %>%
  group_by(variable, MANUAL_CODE, LABEL)

## Warning: attributes are not identical across measure variables;
## they will be dropped

prob_summary%>%
  ggplot(aes(x=as.numeric(as.character(PROB)), fill=factor(MANUAL_CODE)))+
  facet_wrap(~variable)+
  geom_histogram()

## 'stat_bin()' using 'bins = 30'. Pick better
## value with 'binwidth'.
```



```
t(analytics_res@algorithm_summary)
```

```
##           0      1
## SVM_PRECISION 0.98 0.90
## SVM_RECALL    0.79 0.99
## SVM_FSCORE    0.87 0.94
## LOGITBOOST_PRECISION 0.98 0.90
## LOGITBOOST_RECALL 0.79 0.99
## LOGITBOOST_FSCORE 0.87 0.94
## BAGGING_PRECISION 0.98 0.90
## BAGGING_RECALL 0.79 0.99
## BAGGING_FSCORE 0.87 0.94
## FORESTS_PRECISION 1.00 0.90
## FORESTS_RECALL 0.77 1.00
```

```

## FORESTS_FSCORE      0.87 0.95
## TREE_PRECISION      1.00 0.90
## TREE_RECALL         0.78 1.00
## TREE_FSCORE         0.88 0.95
## MAXENTROPY_PRECISION 1.00 0.90
## MAXENTROPY_RECALL   0.78 1.00
## MAXENTROPY_FSCORE   0.88 0.95

t(analytics_res@label_summary)

##
##
## NUM_MANUALLY_CODED      185.00000
## NUM_CONSENSUS_CODED     149.00000
## NUM_PROBABILITY_CODED   149.00000
## PCT_CONSENSUS_CODED     80.54054
## PCT_PROBABILITY_CODED   80.54054
## PCT_CORRECTLY_CODED_CONSENSUS 78.91892
## PCT_CORRECTLY_CODED_PROBABILITY 78.91892
##
##
## NUM_MANUALLY_CODED      362.00000
## NUM_CONSENSUS_CODED     398.00000
## NUM_PROBABILITY_CODED   398.00000
## PCT_CONSENSUS_CODED     109.94475
## PCT_PROBABILITY_CODED   109.94475
## PCT_CORRECTLY_CODED_CONSENSUS 99.17127
## PCT_CORRECTLY_CODED_PROBABILITY 99.17127

analytics_res@ensemble_summary

##          n-ENSEMBLE COVERAGE n-ENSEMBLE RECALL
## n >= 1          1.00          0.92
## n >= 2          1.00          0.92
## n >= 3          1.00          0.92
## n >= 4          0.99          0.93
## n >= 5          0.99          0.93
## n >= 6          0.99          0.93

summary(analytics_res)

## ENSEMBLE SUMMARY
##
##          n-ENSEMBLE COVERAGE n-ENSEMBLE RECALL
## n >= 1          1.00          0.92
## n >= 2          1.00          0.92
## n >= 3          1.00          0.92
## n >= 4          0.99          0.93
## n >= 5          0.99          0.93
## n >= 6          0.99          0.93
##
##
## ALGORITHM PERFORMANCE
##
##          SVM_PRECISION          SVM_RECALL
##          0.940          0.890
##          SVM_FSCORE LOGITBOOST_PRECISION

```

##	0.905	0.940
##	LOGITBOOST_RECALL	LOGITBOOST_FSCORE
##	0.890	0.905
##	BAGGING_PRECISION	BAGGING_RECALL
##	0.940	0.890
##	BAGGING_FSCORE	FORESTS_PRECISION
##	0.905	0.950
##	FORESTS_RECALL	FORESTS_FSCORE
##	0.885	0.910
##	TREE_PRECISION	TREE_RECALL
##	0.950	0.890
##	TREE_FSCORE	MAXENTROPY_PRECISION
##	0.915	0.950
##	MAXENTROPY_RECALL	MAXENTROPY_FSCORE
##	0.890	0.915