

We only have access to a big data set via boolean search.

Sub-sets of the data set returned by a single key word on a single day are large. For example for the single search term ‘election’ for the single day of 1 August 1868 returns 1982 articles. The search term ‘riot’ for the same day (close to the 1868 General Election) returns 360 articles, 228 of which do not also contain the word election.

With even a small number of search terms it quickly becomes impractical to examine all the documents even for a single day.

Following King et al we define S - the search set of all documents in the British Newspaper Archive T - the target set of all documents in the British Newspaper Archive which are about election violence R - a reference set of documents which are about election violence

The task is to identify T from S in a form where T can be

It is trivial to define an algorithm which obtains a subset of S which contains T , because $S \subseteq S$ and $T \subset S$. Algorithms which aim to maximise the chances of obtaining all of T will tend to return S .

Our task is to find a good method for returning T from S in a form which we can analyse. By a *good* method we mean a method which returns a greater proportion of T , and a greater ratio of T to $\neg T$ than alternative methods. The main alternative method is manual searching by historians.

1 The data

```
classdocs <- durhamevp::get_classified_docs()

# just some candidate document for speed/space reasons
unclassdocs <- durhamevp::get_candidate_documents(3000:6000)
```

In the classified data there are 1925 cases:

- **703** non-election articles.
- **796** election violence articles.
- **426** election (but not violent) articles.

2 Keyword Identification

Note it is important to make keyword identification somewhat selective of T from S otherwise even very good stage 2 & 3 selection processes the false positives will overwhelm the true positives.

Algorithm:

1. use classifier on R and S to identify two lists of keywords
2. generate probability from classifier parameters

3. add to keyword list based on probability

```
classified_corpus<-quanteda::corpus(classdocs[,c("fakeid", "ocr", "EV_article")], text_field="ocr")
classified_dfm <- preprocess_corpus(classified_corpus, stem=FALSE, min_termfreq=20, min_docfreq=20)

keywords<-nb_keywords(classified_dfm, "EV_article")
knitr::kable(head(keywords, 20))
```

	rowname	0	1	id
12432	roughly	0.0274481	0.9725519	12432
12433	riotously	0.0310748	0.9689252	12433
12409	smashing	0.0340764	0.9659236	12409
12431	bicycle	0.0358056	0.9641944	12431
12435	pontypool	0.0358056	0.9641944	12435
11726	smashed	0.0476850	0.9523150	11726
12436	bedminster	0.0479798	0.9520202	12436
12408	p.c	0.0544608	0.9455392	12408
12418	roughs	0.0555324	0.9444676	12418
12169	stoned	0.0602765	0.9397235	12169
12429	missile	0.0602765	0.9397235	12429
12434	abersychan	0.0602765	0.9397235	12434
12437	lofts	0.0602765	0.9397235	12437
12386	rioters	0.0641103	0.9358897	12386
9213	staves	0.0752830	0.9247170	9213
11863	supt	0.0766460	0.9233540	11863
12080	foley	0.0766460	0.9233540	12080
12430	polioe	0.0810481	0.9189519	12430
12223	panes	0.0834444	0.9165556	12223
11931	detective	0.0842750	0.9157250	11931

3 Refinement Using Description

```
description_corpus<-quanteda::corpus(classdocs[,c("fakeid", "description", "EV_article")], text_field="description")
description_dfm <- preprocess_corpus(description_corpus, stem=FALSE, min_termfreq=5, min_docfreq=5)

both_dfms<-split_dfm(description_dfm, n_train = 700)

nb<-quanteda::textmodel_nb(both_dfms$train, y=quanteda::docvars(both_dfms$train, "EV_article"))
prob_nb<-predict(nb, newdata = both_dfms$test, type="probability")
pred_nb<-data.frame(predict(nb, newdata = both_dfms$test, type="class"))
res<-data.frame(predict_nb=pred_nb, prob_nb)
names(res)<-c("predict_nb", "prob_notev", "prob_ev")
assess_classification(organise_results(both_dfms$test, res)) %>%
```

```
filter(rowname %in% c("Precision", "Recall", "F1")) %>%
kable()
```

rowname	value	model
Precision	0.7697842	naive bayes
Recall	0.8560000	naive bayes
F1	0.8106061	naive bayes

```
description_corpus<-quantda::corpus(classdocs[,c("fakeid", "description", "EV_article")], t
description_dfm <- preprocess_corpus(description_corpus, stem=FALSE, min_termfreq=5, min_doc

both_dfms<-split_dfm(description_dfm, n_train = 700)
```

Note: this way of creating the dfm does make dfms terms equal because one overall dfm is created and then it is subset. You can see below that the number of features in both dfms is the same:

```
print(both_dfms$train)

## Document-feature matrix of: 700 documents, 1,675 features (99.2% sparse).

print(both_dfms$test)

## Document-feature matrix of: 1,225 documents, 1,675 features (99.2% sparse).
```

3.1 More realistic use case

Exclude non-election cases which would not be returned by our keywords:

```
description_corpus<-quantda::corpus(classdocs[classdocs$selection_article==1,c("fakeid", "d
description_dfm <- preprocess_corpus(description_corpus, stem=FALSE, min_termfreq=5, min_doc

both_dfms<-split_dfm(description_dfm, n_train = 700)

nb<-quantda::textmodel_nb(both_dfms$train, y=quantda::docvars(both_dfms$train, "EV_article
prob_nb<-predict(nb, newdata = both_dfms$test, type="probability")
pred_nb<-data.frame(predict(nb, newdata = both_dfms$test, type="class"))
res<-data.frame(predict_nb=pred_nb, prob_nb)
names(res)<-c("predict_nb", "prob_notev", "prob_ev")
assess_classification(organise_results(both_dfms$test, res)) %>%
  filter(rowname %in% c("Precision", "Recall", "F1")) %>%
  kable()
```

rowname	value	model
Precision	0.8750000	naive bayes
Recall	0.7754491	naive bayes
F1	0.8222222	naive bayes

```
description_corpus<-quantda::corpus(classdocs[,c("fakeid", "description", "EV_article")], t
description_dfm <- preprocess_corpus(description_corpus, stem=FALSE, min_termfreq=5, min_docf
both_dfms<-split_dfm(description_dfm, n_train = 1000)
```

4 Refinement Using OCR

```
ocr_corpus<-quantda::corpus(classdocs[,c("fakeid", "ocr", "EV_article")], text_field = "ocr"
ocr_dfm <- preprocess_corpus(ocr_corpus, stem=FALSE, min_termfreq=5, min_docfreq = 2)

both_dfms<-split_dfm(ocr_dfm, n_train = 700)
nb<-quantda::textmodel_nb(both_dfms$train, y=quantda::docvars(both_dfms$train, "EV_article")
prob_nb<-predict(nb, newdata = both_dfms$test, type="probability")
pred_nb<-data.frame(predict(nb, newdata = both_dfms$test, type="class"))
res<-data.frame(predict_nb=pred_nb, prob_nb)
names(res)<-c("predict_nb", "prob_notev", "prob_ev")
assess_classification(organise_results(both_dfms$test, res)) %>%
  filter(rowname %in% c("Precision", "Recall", "F1")) %>%
  kable()
```

rowname	value	model
Precision	0.6708683	naive bayes
Recall	0.9373777	naive bayes
F1	0.7820408	naive bayes

5 King Algorithm: implementation

5.1 Incrementally Define R and S

R is our reference set. [King suggestions: could define R based on one simple keyword search].

5.1.1 Intermediate step

Take keywords in R , K_R , ranked by simple statistic such as document frequency or frequency-inverse document frequency. User examines elements of K_R apart from those used to define the set and chooses some keywords to define Q_S , which in turn generates a definition for S so that we can run the rest of the algorithm. The user can continue to add keywords from K_R into the final desired query Q_{RT} . This step also mitigates the issue of how to define a search set in large data sets that do not fit into memory all at once or may not even be able to be retrieved all at once. is the BNA.

5.2 Partition S into T and $S \setminus T$

To partition S into T and $S \setminus T$, we first define a ‘training’ set by sampling from S and R . Since R is typically much smaller than S our test set for our classifiers is all all of S , we often use the entire R set and a sample of S as our training set.

```
R <- classdocs[classdocs$EV_article==1,]
R$R <- 1
R$in_sample<-1
S <- unclassdocs
S$R <- 0

n_sample_S <- 700
S$in_sample<-0
S$in_sample[sample(1:nrow(S), n_sample_S)]<-1

R_S <- dplyr::bind_rows(R, S)
R_S$fakeid<-1:nrow(R_S)

R_S_corpus<-quanteda::corpus(R_S[,c("fakeid", "ocr", "in_sample", "R")], text_field = "ocr")
R_S_dfm <- preprocess_corpus(R_S_corpus, stem=FALSE, min_termfreq=20, min_docfreq = 20)

king_train_dfm<-quanteda::dfm_subset(R_S_dfm, quanteda::docvars(R_S_dfm, "in_sample")==1)
king_nb <- quanteda::textmodel_nb(king_train_dfm, y=quanteda::docvars(king_train_dfm, "R"),
keywords<-nb_keywords(king_train_dfm, "R")
knitr::kable(head(keywords, 20))
```

	rowname	0	1	id
3	generated	0.0323522	0.9676478	3
14	legible	0.0332999	0.9667001	14
2921	chartist	0.0343049	0.9656951	2921
5	optical	0.0365084	0.9634916	5
8	technology	0.0365084	0.9634916	8
13	deciphering	0.0365084	0.9634916	13
2501	p.c	0.0443407	0.9556593	2501
2671	detective	0.0470963	0.9529037	2671
13000	guttled	0.0513514	0.9486486	13000
9243	bicycle	0.0564516	0.9435484	9243
2271	roughs	0.0594015	0.9405985	2271
1079	cobden	0.0626767	0.9373233	1079
9441	supt	0.0626767	0.9373233	9441
2952	conser	0.0663342	0.9336658	2952
3425	smashing	0.0704449	0.9295551	3425
3470	inciting	0.0804111	0.9195889	3470
9449	schoolroom	0.0804111	0.9195889	9449
10758	lodgers	0.0804111	0.9195889	10758
13297	gladstone's	0.0804111	0.9195889	13297
1995	shutters	0.0865322	0.9134678	1995

After fitting the classifiers, we use the estimated parameters to generate predicted probabilities of R membership for all documents in S . Of course, all the search set documents in fact fall within S but our interest is in learning from the *mistakes* these classifiers make.

```
S_dfm <- quanteda::dfm_subset(R_S_dfm, quanteda::docvars(R_S_dfm, "R")==0)
quanteda::docvars(S_dfm, "T")<-predict(king_nb, newdata = S_dfm, type="class")
```

5.3 Discovering Keywords to Classify Documents

After partitioning S into estimated target T and non-target set $S \setminus T$, we must find and rank keywords which best discriminate T and $S \setminus T$. King does this in three stages:

5.3.1 identify keywords in S

5.3.2 sort them into those that predict each of the two steps

My suggestion here is simply to use binary_dfm

```
S_dfm_binary<-quanteda::dfm_weight(S_dfm, scheme="boolean")
king_stage2 <- quanteda::textmodel_nb(S_dfm, y=quanteda::docvars(S_dfm_binary, "T"), distrib="binomial")
king_stage2 <- nb_keywords(S_dfm, "T")
knitr::kable(head(king_stage2, 30))
```

	rowname	0	1	id
8565	horden	0.0152274	0.9847726	8565
1208	protectionist	0.0186758	0.9813242	1208
2335	dump	0.0210604	0.9789396	2335
9323	suffragists	0.0219967	0.9780033	9323
8621	motor	0.0256448	0.9743552	8621
9243	bicycle	0.0341371	0.9658629	9243
9137	truncheons	0.0430465	0.9569535	9137
9404	rowdyism	0.0471482	0.9528518	9404
9449	schoolroom	0.0471482	0.9528518	9449
9530	telegraphed	0.0471482	0.9528518	9530
9120	batons	0.0490164	0.9509836	9120
2417	workers	0.0507769	0.9492231	2417
9702	nationalist	0.0521138	0.9478862	9702
9261	railings	0.0540100	0.9459900	9261
9325	barricaded	0.0540100	0.9459900	9325
8852	unionism	0.0550107	0.9449893	8852
2921	chartist	0.0582486	0.9417514	2921
9168	booing	0.0582486	0.9417514	9168
8742	drafted	0.0582486	0.9417514	8742
8556	nationalists	0.0602760	0.9397240	8556
9966	socialist	0.0602760	0.9397240	9966
8867	industries	0.0611270	0.9388730	8867
8534	hartlepool	0.0618915	0.9381085	8534
8876	candidature	0.0618915	0.9381085	8876
8566	sacked	0.0643045	0.9356955	8566
3425	smashing	0.0660205	0.9339795	3425
2632	unionist	0.0673033	0.9326967	2632
10244	socialism	0.0690934	0.9309066	10244
3076	asquith	0.0742795	0.9257205	3076
1993	witness's	0.0761856	0.9238144	1993

```
knitr::kable(tail(king_stage2, 30))
```

	rowname	0	1	id
10442	mankind	0.9348514	0.0651486	10442
12687	essence	0.9348514	0.0651486	12687
16592	regency	0.9348514	0.0651486	16592
10466	relaxed	0.9358848	0.0641152	10466
13348	il	0.9358848	0.0641152	13348
16565	berchem	0.9358848	0.0641152	16565
13432	spoliation	0.9368860	0.0631140	13432
4914	forts	0.9378563	0.0621437	4914
11212	stupid	0.9378563	0.0621437	11212
16087	talleyrand	0.9378563	0.0621437	16087
16374	honored	0.9387973	0.0612027	16374
7535	batteries	0.9392572	0.0607428	7535
14632	wvill	0.9397102	0.0602898	14632
14892	unfettered	0.9397102	0.0602898	14892
13005	prussia	0.9403038	0.0596962	13005
16672	despised	0.9405963	0.0594037	16672
8124	scheldt	0.9418776	0.0581224	8124
6584	belgium	0.9427826	0.0572174	6584
8113	lunette	0.9435024	0.0564976	8113
6666	ambition	0.9436337	0.0563663	6666
14407	belgian	0.9455937	0.0544063	14407
8114	laurent	0.9468458	0.0531542	8114
14963	engraved	0.9482080	0.0517920	14963
15385	carolina	0.9488633	0.0511367	15385
6939	combat	0.9501253	0.0498747	6939
4130	innovation	0.9579148	0.0420852	4130
11224	senate	0.9595975	0.0404025	11224
15979	frontiers	0.9607738	0.0392262	15979
8121	bombs	0.9696190	0.0303810	8121
8117	besiegers	0.9747445	0.0252555	8117

5.4 rank them by degree of discriminatory power

We already seem to have done this in the step above

5.5 Human Input and Human-Computer Iteration

The King et al suggestion is to present the above lists to humans who then choose additional keywords from these lists. That might be fine if the final goal is develop keywords. However, this does not work so well if the final goal is to obtain sets of documents based on those keywords. Some problems with this approach for obtaining documents

1. the set of documents will keep on growing without limit (with many irrelevant documents)

2. there is no guidance about which keywords to select from the lists presented

An alternative proposal is to continue with a classification approach. That is to treat the classifier as informative about two separate matters: keywords retrieval as potentially informative

The stage II classifier is informative about the following issues:

1. terms which are in S but not in R which may be positively predictive of election violence.
2. terms which are very differently predictive of election violence between stage I and stage II classification.
- 3.

Idea 1: For keyword t in keywords $1 \dots n$, Q_t is the Boolean query which returns document d from S if d contains t . S' is the set of documents returned by the set of queries $Q_{T_1} \dots Q_{T_n}$. The probability that document d is an example of the concept of interest is given by the set of queries which return that document. We run independent queries with each of the keywords. We have different sets returned from

Idea 2: King's stage II classification can be thought of as the first part of an EM process which augments information from labelled data with information from unlabelled data. See for a classic example Nigam et. al. 2000 - Semi-supervised parameter estimation: can learn from a combination of labeled and unlabeled data in a loop. First train a classifier using available labeled documents, and probabilistically label the unlabeled documents. Then train a new classifier using the labels for all the documents, iterate to convergence. This basic EM procedure works well when the data conform to the generative assumptions of the model. However, these assumptions are often violated in practice, and poor performance can result. Two extensions improve classification accuracy (1) a weighting factor to modulate the contribution of the unlabeled data and (2) the use of multiple mixture components per class.

How can unlabeled data increase classification accuracy? Unlabelled data provide information about the joint probability distribution over words. Suppose that using only the labeled data we determine that documents containing the word 'riot' belongs to the election violence class. If we use this fact to estimate the classification of the many unlabeled documents, we might find the word 'bludgeon' occurs frequently with the word 'riot' and we might construct a more accurate classifier that considers both 'riot' and 'bludgeon' as indicators of positive examples.

Given a training set $(x^{(i)}, y^{(i)})$ for $i = 1 \dots n$, where θ is a parameter vector consisting of the values for all parameters $q(y)$ and $q_i(x|y)$ the log-likelihood

function is:

$$\begin{aligned}
 L(\theta) &= \sum_{i=1}^n \log p(x^{(i)}, y^{(i)}) \\
 &= \sum_{i=1}^n \log q(y^{(i)}) + \sum_{i=1}^n \sum_{j=1}^d \log q_j(x_j^{(i)} | y^{(i)})
 \end{aligned} \tag{1}$$

Idea 3: to find a set of keywords searches which returns T - when you have a set of keywords which returns T - adding new keywords will change S but not reveal more cases of T - so will not see large changes in individual word predictivity of class in Naive Bayes classifier. encompasses the add search terms until revised data set S' is not substantially changed by addition of new When new search terms are added this changes the nature of S - at each iteration new information added - consider changes to S'

```

names(king_stage2)<-c("rowname", "stage2_0", "stage2_1", "stage2_id")
names(keywords)<-c("rowname", "stage1_0", "stage1_1", "stage1_id")
left_join(king_stage2, keywords, by="rowname") %>%
  mutate(logit_stage1=log(stage1_1/stage1_0), logit_stage2=log(stage2_1/stage2_0)) %>%
  ggplot(aes(logit_stage1, logit_stage2))+
  geom_point()

```

